

# ml lab homework 9

March 7, 2024

## 1 Machine Learning Lab Homework 9

2 Connor O’Keefe

3 03/07/2024

## 4 0.) Import and Clean data

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

```
[2]: df = pd.read_csv("Country-data.csv", sep = ",")
```

```
[3]: df
```

```
[3]:
```

	country	child_mort	exports	health	imports	income	\
0	Afghanistan	90.2	10.0	7.58	44.9	1610	
1	Albania	16.6	28.0	6.55	48.6	9930	
2	Algeria	27.3	38.4	4.17	31.4	12900	
3	Angola	119.0	62.3	2.85	42.9	5900	
4	Antigua and Barbuda	10.3	45.5	6.03	58.9	19100	
..	...	...	...	...	...		
162	Vanuatu	29.2	46.6	5.25	52.7	2950	
163	Venezuela	17.1	28.5	4.91	17.6	16500	
164	Vietnam	23.3	72.0	6.84	80.2	4490	
165	Yemen	56.3	30.0	5.18	34.4	4480	
166	Zambia	83.1	37.0	5.89	30.9	3280	

	inflation	life_expec	total_fer	gdpp
0	9.44	56.2	5.82	553
1	4.49	76.3	1.65	4090
2	16.10	76.5	2.89	4460
3	22.40	60.1	6.16	3530

4	1.44	76.8	2.13	12200
..	...	...	...	...
162	2.62	63.0	3.50	2970
163	45.90	75.4	2.47	13500
164	12.10	73.1	1.95	1310
165	23.60	67.5	4.67	1310
166	14.00	52.0	5.40	1460

[167 rows x 10 columns]

```
[4]: X = df.drop('country', axis=1)
      scaler = StandardScaler().fit(X)
      X_scaled = scaler.transform(X)
```

Question we want to answer: Can k-means identify developing economies?

## 5 1.) Fit a kmeans Model with any Number of Clusters

```
[5]: kmeans = KMeans(n_clusters = 5).fit(X_scaled)
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:  
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
there are less chunks than available threads. You can avoid it by setting the  
environment variable OMP\_NUM\_THREADS=1.  
warnings.warn(

## 6 2.) Pick two features to visualize across

```
[6]: X.columns
```

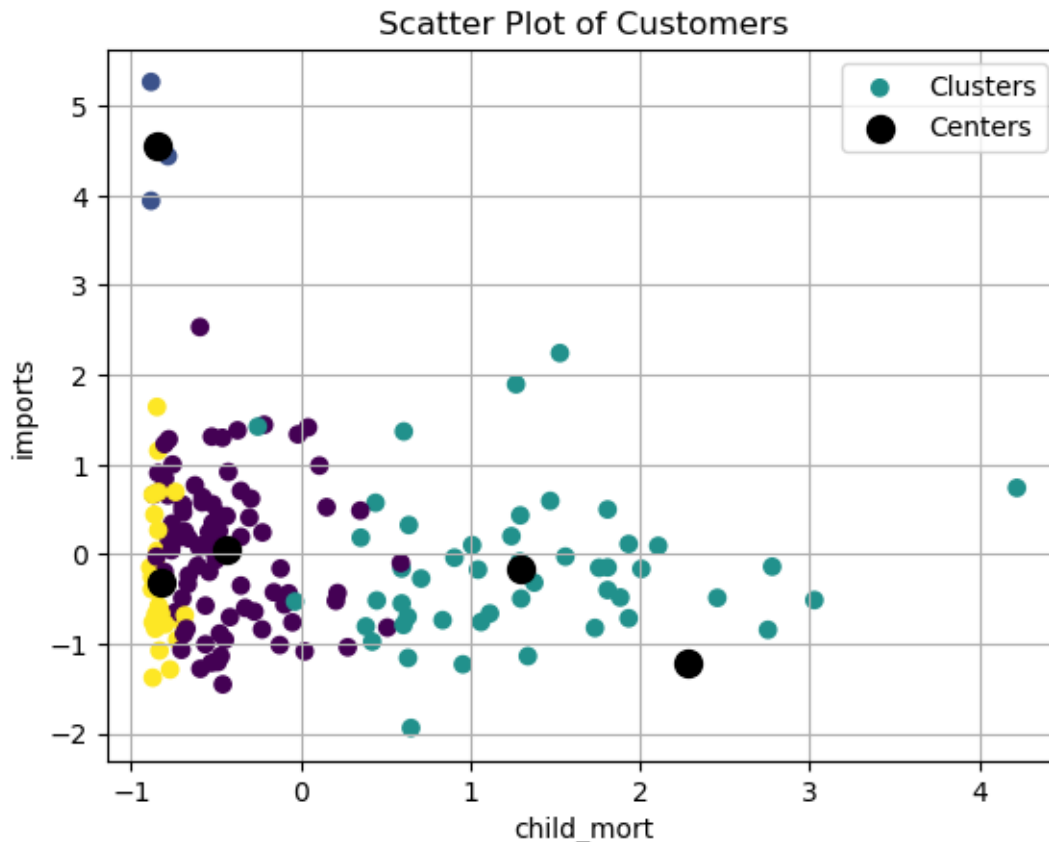
```
[6]: Index(['child_mort', 'exports', 'health', 'imports', 'income', 'inflation',
          'life_expec', 'total_fer', 'gdpp'],
          dtype='object')
```

```
[7]: # picking what variables to compare below
      # child mortality vs imports
      x1_index = 0
      x2_index = 3

      scatter = plt.scatter(X_scaled[:, x1_index], X_scaled[:, x2_index], c=kmeans.
          ↪labels_, cmap='viridis', label='Clusters')

      centers = plt.scatter(kmeans.cluster_centers_[:, x1_index], kmeans.
          ↪cluster_centers_[:, x2_index], marker='o', color='black', s=100,
          ↪label='Centers')
```

```
plt.xlabel(X.columns[x1_index])
plt.ylabel(X.columns[x2_index])
plt.title('Scatter Plot of Customers')
plt.legend()
plt.grid()
plt.show()
```



- 7 3.) Check a range of k-clusters and visualize to find the elbow.  
Test 30 different random starting places for the centroid means

```
[8]: WCSSs = []
Ks = range(1,15)
for k in Ks:
    kmeans = KMeans(n_clusters = k, n_init = 30).fit(X_scaled)
    WCSSs.append(kmeans.inertia_)
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:  
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
there are less chunks than available threads. You can avoid it by setting the



UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:

UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:

UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```


C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:

UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```

**8 4.) Use the above work and economic critical thinking to choose a number of clusters. Explain why you chose the number of clusters and fit a model accordingly.**

```
[10]: plt.plot(WCSSs)
plt.xlabel('# of clusters')
plt.ylabel('WCSS')
plt.show()
```



As can be seen above, there is no real elbow. This does not mean that the method is not working, just that the ideal number of clusters is up for interpretation and dependent on the context. Choosing two clusters for comparison seems to be an ideal number, because we could compare affluent countries to impoverished countries.

## 9 5.) Do the same for a silhouette plot

```
[11]: from sklearn.metrics import silhouette_score
```

```
[12]: SSs = []
Ks = range(2,15) # can't only have one cluster because no other cluster to
    ↪ compare with
for k in Ks:
    kmeans = KMeans(n_clusters = k, n_init = 30).fit(X_scaled)
    sil = silhouette_score(X_scaled, kmeans.labels_)
    SSs.append(sil)
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:  
 UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
 there are less chunks than available threads. You can avoid it by setting the  
 environment variable OMP\_NUM\_THREADS=1.

```

warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.

```

there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.

```
warnings.warn(
```

```
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
```

```
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
there are less chunks than available threads. You can avoid it by setting the  
environment variable OMP_NUM_THREADS=1.
```


```
warnings.warn(
```

```
C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1446:
```

```
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
there are less chunks than available threads. You can avoid it by setting the  
environment variable OMP_NUM_THREADS=1.
```

```
warnings.warn(
```

```
[13]: plt.plot(Ks, SSs)  
plt.xlabel('# of clusters')  
plt.ylabel('WCSS')  
plt.show()
```





## 10 6.) Create a list of the countries that are in each cluster. Write interesting things you notice.

```
[14]: # choosing two clusters below
kmeans = KMeans(n_clusters = 2, n_init = 30).fit(X_scaled)
```

C:\Users\12282\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1446:  
 UserWarning: KMeans is known to have a memory leak on Windows with MKL, when  
 there are less chunks than available threads. You can avoid it by setting the  
 environment variable OMP\_NUM\_THREADS=1.  
 warnings.warn(

```
[15]: preds = pd.DataFrame(kmeans.labels_)
```

```
[16]: output = pd.concat([preds, df], axis = 1)
output
```

```
[16]:      0      country  child_mort  exports  health  imports  income  \
0  0      Afghanistan      90.2      10.0      7.58      44.9      1610
1  1          Albania      16.6      28.0      6.55      48.6      9930
2  1          Algeria      27.3      38.4      4.17      31.4     12900
3  0          Angola     119.0      62.3      2.85      42.9      5900
4  1  Antigua and Barbuda      10.3      45.5      6.03      58.9     19100
..  ..          ...          ...          ...          ...          ...
162 0          Vanuatu      29.2      46.6      5.25      52.7      2950
163 1          Venezuela      17.1      28.5      4.91      17.6     16500
164 1          Vietnam      23.3      72.0      6.84      80.2      4490
165 0          Yemen       56.3      30.0      5.18      34.4      4480
166 0          Zambia      83.1      37.0      5.89      30.9      3280

      inflation  life_expec  total_fer  gdpp
0           9.44       56.2       5.82   553
1           4.49       76.3       1.65  4090
2          16.10       76.5       2.89  4460
3          22.40       60.1       6.16  3530
4           1.44       76.8       2.13 12200
..          ...          ...          ...   ...
162          2.62       63.0       3.50  2970
163         45.90       75.4       2.47 13500
164         12.10       73.1       1.95  1310
165         23.60       67.5       4.67  1310
166         14.00       52.0       5.40  1460
```

[167 rows x 11 columns]

```
[17]: cluster1 = list(output.loc[output[0] == 0, 'country'])
print('Cluster 1: ', cluster1)
```

```
Cluster 1: ['Afghanistan', 'Angola', 'Bangladesh', 'Benin', 'Bolivia',
'Botswana', 'Burkina Faso', 'Burundi', 'Cambodia', 'Cameroon', 'Central African
Republic', 'Chad', 'Comoros', 'Congo, Dem. Rep.', 'Congo, Rep.', 'Cote
d'Ivoire', 'Egypt', 'Equatorial Guinea', 'Eritrea', 'Gabon', 'Gambia', 'Ghana',
'Guatemala', 'Guinea', 'Guinea-Bissau', 'Guyana', 'Haiti', 'India', 'Indonesia',
'Iraq', 'Kenya', 'Kiribati', 'Kyrgyz Republic', 'Lao', 'Lesotho', 'Liberia',
'Madagascar', 'Malawi', 'Mali', 'Mauritania', 'Micronesia, Fed. Sts.',
'Mongolia', 'Mozambique', 'Myanmar', 'Namibia', 'Nepal', 'Niger', 'Nigeria',
'Pakistan', 'Philippines', 'Rwanda', 'Samoa', 'Senegal', 'Sierra Leone',
'Solomon Islands', 'South Africa', 'Sudan', 'Tajikistan', 'Tanzania', 'Timor-
Leste', 'Togo', 'Tonga', 'Turkmenistan', 'Uganda', 'Uzbekistan', 'Vanuatu',
'Yemen', 'Zambia']
```

```
[18]: cluster2 = list(output.loc[output[0] == 1, 'country'])
print('Cluster 2: ', cluster2)
```

```
Cluster 2: ['Albania', 'Algeria', 'Antigua and Barbuda', 'Argentina',
'Armenia', 'Australia', 'Austria', 'Azerbaijan', 'Bahamas', 'Bahrain',
'Barbados', 'Belarus', 'Belgium', 'Belize', 'Bhutan', 'Bosnia and Herzegovina',
'Brazil', 'Brunei', 'Bulgaria', 'Canada', 'Cape Verde', 'Chile', 'China',
'Colombia', 'Costa Rica', 'Croatia', 'Cyprus', 'Czech Republic', 'Denmark',
'Dominican Republic', 'Ecuador', 'El Salvador', 'Estonia', 'Fiji', 'Finland',
'France', 'Georgia', 'Germany', 'Greece', 'Grenada', 'Hungary', 'Iceland',
'Iran', 'Ireland', 'Israel', 'Italy', 'Jamaica', 'Japan', 'Jordan',
'Kazakhstan', 'Kuwait', 'Latvia', 'Lebanon', 'Libya', 'Lithuania', 'Luxembourg',
'Macedonia, FYR', 'Malaysia', 'Maldives', 'Malta', 'Mauritius', 'Moldova',
'Montenegro', 'Morocco', 'Netherlands', 'New Zealand', 'Norway', 'Oman',
'Panama', 'Paraguay', 'Peru', 'Poland', 'Portugal', 'Qatar', 'Romania',
'Russia', 'Saudi Arabia', 'Serbia', 'Seychelles', 'Singapore', 'Slovak
Republic', 'Slovenia', 'South Korea', 'Spain', 'Sri Lanka', 'St. Vincent and the
Grenadines', 'Suriname', 'Sweden', 'Switzerland', 'Thailand', 'Tunisia',
'Turkey', 'Ukraine', 'United Arab Emirates', 'United Kingdom', 'United States',
'Uruguay', 'Venezuela', 'Vietnam']
```

Because I am an American, my first point of interest was indentifying how the United States was classified (Cluster 2). Overall, Cluster 2 seems to classify the more developed countries and Cluster 1 seems to classify developing countries, with a fair amount from Africa and Asia. More analysis follows below.

## 11 7.) Create a table of Descriptive Statistics. Rows being the Cluster number and columns being all the features. Values being the mean of the centroid. Use the nonscaled X values for interprotation

```
[21]: output.drop('country', axis = 1).groupby(0).mean()
```

```
[21]: child_mort    exports    health    imports          income  inflation  \
0
0   76.280882   30.198515   6.090147   43.642146    4227.397059   11.098750
1   12.161616   48.603030   7.314040   49.121212   26017.171717    5.503545

    life_expec  total_fer          gdpp
0
0   61.910294    4.413824    1981.235294
1   76.493939    1.941111   20507.979798
```

```
[20]: output.drop('country', axis = 1).groupby(0).std()
```

```
[20]: child_mort    exports    health    imports          income  inflation  \
0
0   38.076068   18.201742   2.645319   19.323451    4890.581414   13.682630
1    8.523122   30.116032   2.716652   26.928785   20441.749847    6.957187

    life_expec  total_fer          gdpp
0
0    6.897418    1.285590    2528.509189
1    3.735757    0.486744   20578.727127
```

## 12 8.) Write an observation about the descriptive statistics.

As I hypothesized above, I think Cluster 1 may classify developing countries and Cluster 2 may classify developed countries. So now, turning to the above statistics, some important differences can be noticed. Every “good” and “bad” mean statistic is more optimal for Cluster 2, which supports my claim. For example, mean income of Cluster 2 is more than 6x that of Cluster 1. The standard deviation statistics are a little bit harder to interpret, and can be for a few different reasons. For example, the standard deviation of inflation of the developing countries could be higher because there is a huge variation in inflation rates, impacted by countries like Sri Lanka and Venezuela.